Potentials MCP Connector Guide

Oct 23, 2026 · updated Nov 7, 2026 · @Steen M

1. Installing the connector

PotentialsRepos is a remote MCP server, added once as a custom connector in claude.ai. After that it works in the browser, Claude Desktop and the mobile apps.

You need:

- A Claude plan that allows custom connectors (Free plans are limited to one)
- The Potentials server URL: **https://mcp.innovia.dk/mcp**
- A login for the Potentials server, since the connector asks you to sign in

Steps:

1. Open claude.ai and go to **Customize → Connectors** (claude.ai/customize/connectors).
2. Click **+** and choose **Add custom connector**.
3. Enter the name `PotentialsRepos` and paste the server URL.
4. Click **Add**, then **Connect**, and sign in when asked.
5. In a chat, open the tools menu and check that PotentialsRepos is switched on.
6. Test it by asking: "List the Potentials datasets."

On Team and Enterprise plans, an owner adds the connector first under **Organization settings → Connectors**. In Claude Code, add it from the terminal with `claude mcp add --transport http potentials https://mcp.innovia.dk/mcp`.

Claude reaches the server from Anthropic's cloud, not from your own PC, so the server must be reachable over the public internet. (source)

2. How to use the connector

The Potentials connector gives Claude direct, read-only access to the data in the Potentials repository. Claude can look things up, compute, and draw charts from it without anyone exporting CSVs.

The big idea: what an MCP connector is

MCP (Model Context Protocol) is a standard "plug" between Claude and an outside system. The Potentials server publishes a small menu of **tools** (actions Claude can call) and the data behind them. When you ask a question, Claude picks the right tool, calls it, gets numbers back, and reasons over them. You never have to say which tool to use. Just ask in plain words.

The eight tools

Tool	What it does	Example question
list_datasets	Shows the catalogue of all datasets	"What data do you have?"

Tool	What it does	Example question
search_tickers	Finds tickers by name, sector, GICS group, country or zone	"Find Danish healthcare stocks"
get_series	Gets a time series per ticker (daily, weekly or monthly; raw values or % change). Large results also come as a CSV download (section 3)	"Monthly gains for MSFT and MU this year"
plot_series	Same as get_series, but returns a chart (up to 8 tickers)	"Plot RSI for NOVO-B.CO vs LLY since June"
get_table	Gets rows from table-style datasets. Large results also come as a CSV download (section 3)	"Show today's StrategicStocks picks"
tunnel_report	Builds a trend-channel ("tunnel") report for 1–200 tickers: CSV, PDF, PNGs and download links	"Tunnel report for my top 20 by rank"
tunnel_history	Searches the full tunnel history (every trading day for every ticker) with filters or summaries. Big matches come as a CSV download	"Which tickers had a tunnel of at least 60 days with the close near the upper line last week?"
create_api_key	Creates your personal key for CSV downloads (shown once)	"Create my Potentials API key"

Two kinds of data: matrix vs table

Understanding this distinction is key.

- **Matrix** = one time series per ticker (dates × tickers). Most of these are named `longi_*`. They cover about 2.7 years of history. Read them with get_series or plot_series.
- **Table** = ordinary rows, one per ticker (or several per ticker for history). Read them with get_table.

Matrix families (about 80 datasets)

- **Prices and averages:** `PotDat` (the root price table), `longi_price`, MA10/20/50/200, and price-divided-by-MA ratios
- **Trailing gains, the "seven-pack":** `longi_per1d` up to `longi_per200d`
- **Ranking:** `longi_rank` (lower means stronger), its rolling medians, and the step-up counts (3 means a clean uptrend)
- **Momentum:** RSI14, MACD(4,15,9) with zero-crossing flags called ZOP and ZNEd, and MA quotients
- **Trend quality:** `regr*` (trend gain from a log-price fit) and `regrfit*` (R^2, which shows how *straight* the trend is)
- **Risk:** volatility, Sharpe ratio, beta against the core index and the sector, and spread to the 100- or 250-day high
- **Sector level:** `longi_grp_GICS_*` and `longi_grp_Sector2_*`. In these, the rows are sector names, not tickers.

- **Event rebases:** `longi_trump` (prices set to 1.0 at the 2 Apr 2025 tariff announcement) and `longi_iran` (prices set to 1.0 at war start 27 Feb 2026)

Tables

- **Stamdata:** master data — the system-wide ticker list with sector, GICS, homeland and core index
- **potrank2:** similar the current PotRank Google Sheet
- **StrategicStocks:** today's strategy picks
- **across:** every factor for every ticker on the newest day, in one wide table
- **Yfinance** and **StockData2_stacked:** analyst targets and fundamentals history

⚠️ Important: the longi_future_* datasets

These hold **forward** returns and drawdowns, meaning what happened *after* each day. They are excellent for testing a signal: "When RSI was below 30, what was the average 20-day gain afterwards?"

They must **never** be used as an input for predicting that same day. Doing so is called **look-ahead bias** (or data leakage). It makes a strategy look brilliant in a backtest, but the result is fake because the strategy is secretly peeking at the future. This matters especially for ML and channel-scanner work.

Good first things to try

- "Which GICS sectors are strongest over 20 and 50 days right now?"
- "Show me the 10 best-ranked US tickers with a clean step-up of 3"
- "Did stocks with RSI under 30 outperform over the next 20 days?" (uses the future datasets correctly)
- "Run a tunnel report on today's StrategicStocks"
- "Give me the complete potrank2 table as a CSV download"

3. Downloading complete results (CSV)

Claude shows small answers in the chat. When a result is bigger than fits (many tickers over a long period, a whole table, a long list of tunnel matches), it shows the newest part and also gives you a download link to the complete result as a CSV file. You can ask for the file directly: "Give me the complete potrank2 table as a CSV download." The chat then shows only a 10-row preview and the link, which suits scripts and agents.

Your personal API key

A download link opens only together with your personal key, so a link that is forwarded, or kept by someone who has left the network, opens nothing.

- **Create it:** ask Claude "Create my Potentials API key". The key is shown once. Save it at once in a password manager or an environment variable.
- **Lost or leaked:** ask for a new one. The old key stops working immediately.
- **Do not share it:** every download is logged under your name.
- **When your access ends:** the key stops working with it.

Fetching the file from a script

Send the key as a header with the request (Authorization: Bearer <your key>, or the header X-API-Key). From a terminal:

```
curl -H "Authorization: Bearer $POTENTIALS_KEY" -O "<download link from Claude>"
```

From Python:

```
import requests
r = requests.get(url, headers={"Authorization": f"Bearer {key}"})
open("result.csv", "wb").write(r.content)
```

Good to know

- The file is a standard CSV (comma separator, decimal point), unlike the European format of the repository files. A time series has one row per date and one column per ticker.
- A link lasts 24 hours; ask Claude again if it has expired.
- Clicking the link in a browser shows an error, because a browser cannot send the key. Use a script, curl or any tool that can set a header.
- Tunnel report links (the PDF and PNG files from tunnel_report) are different: they open for anyone who holds the link, for 7 days, so do not forward them.